

OS Support for Building Distributed Applications: Multithreaded Programming using Java Threads



Object-Oriented
Programming
with **JAVA**
Essentials and Applications

Rajkumar Buyya | S Thamarai Selvi | Xingchen Chu

Copyrighted Material

Dr. Rajkumar Buyya

Cloud Computing and **D**istributed **S**ystems (CLOUDS) Laboratory

School of Computing and Information Systems

The University of Melbourne, Australia

<http://www.buyya.com>

Agenda

- Introduction
- Thread Applications
- Defining Threads
- Java Threads and States
- Examples

Introduction

- In a networked world, it is common practice to share resources among multiple users.
- Even on desktop computers, users typically run multiple applications and carry out some operations in the background (e.g., printing) and some in the foreground (e.g., editing) simultaneously.
- Modern programming languages and operating systems are designed to support the development of applications containing multiple activities that can be executed concurrently
- **Preemptive** (processor responsibility) **vs. cooperative** (process responsibility) **multitasking?**
- **Since the processor is switching between the applications at intervals of milliseconds, you feel that all applications run concurrently**

Threaded Applications

■ Modern Systems

- Multiple applications run concurrently!
- This means that... there are multiple processes on your computer



A single threaded program

```
class ABC
```

```
{
```

```
....
```

```
    public void main(..)
```

```
    {
```

```
    ...
```

```
    ..
```

```
    }
```

```
}
```

begin

body

end



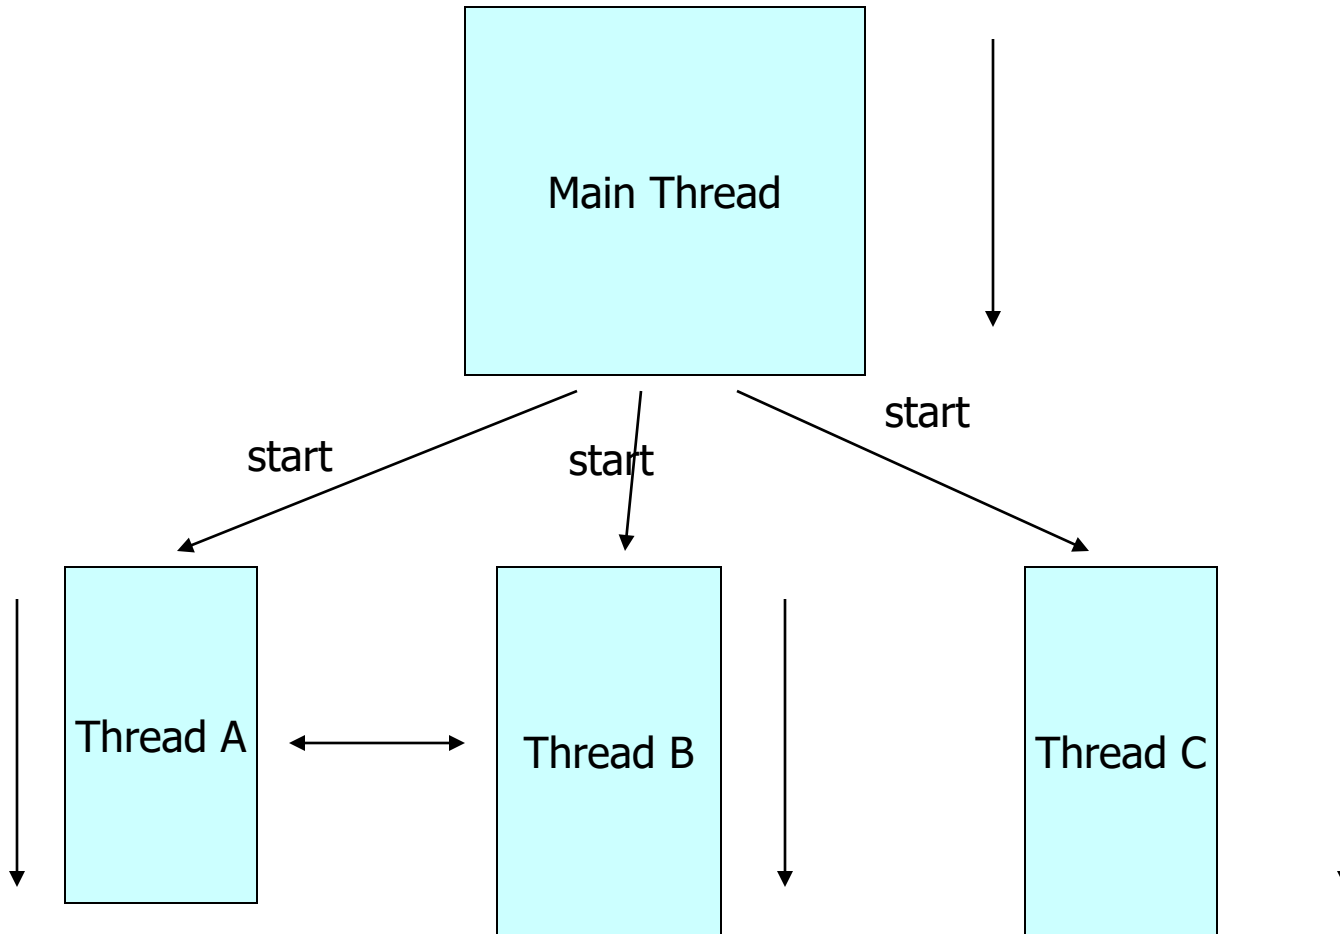
Threaded Applications

■ Modern Systems

- Applications perform many tasks at once!
- This means that... there are multiple threads within a single process.



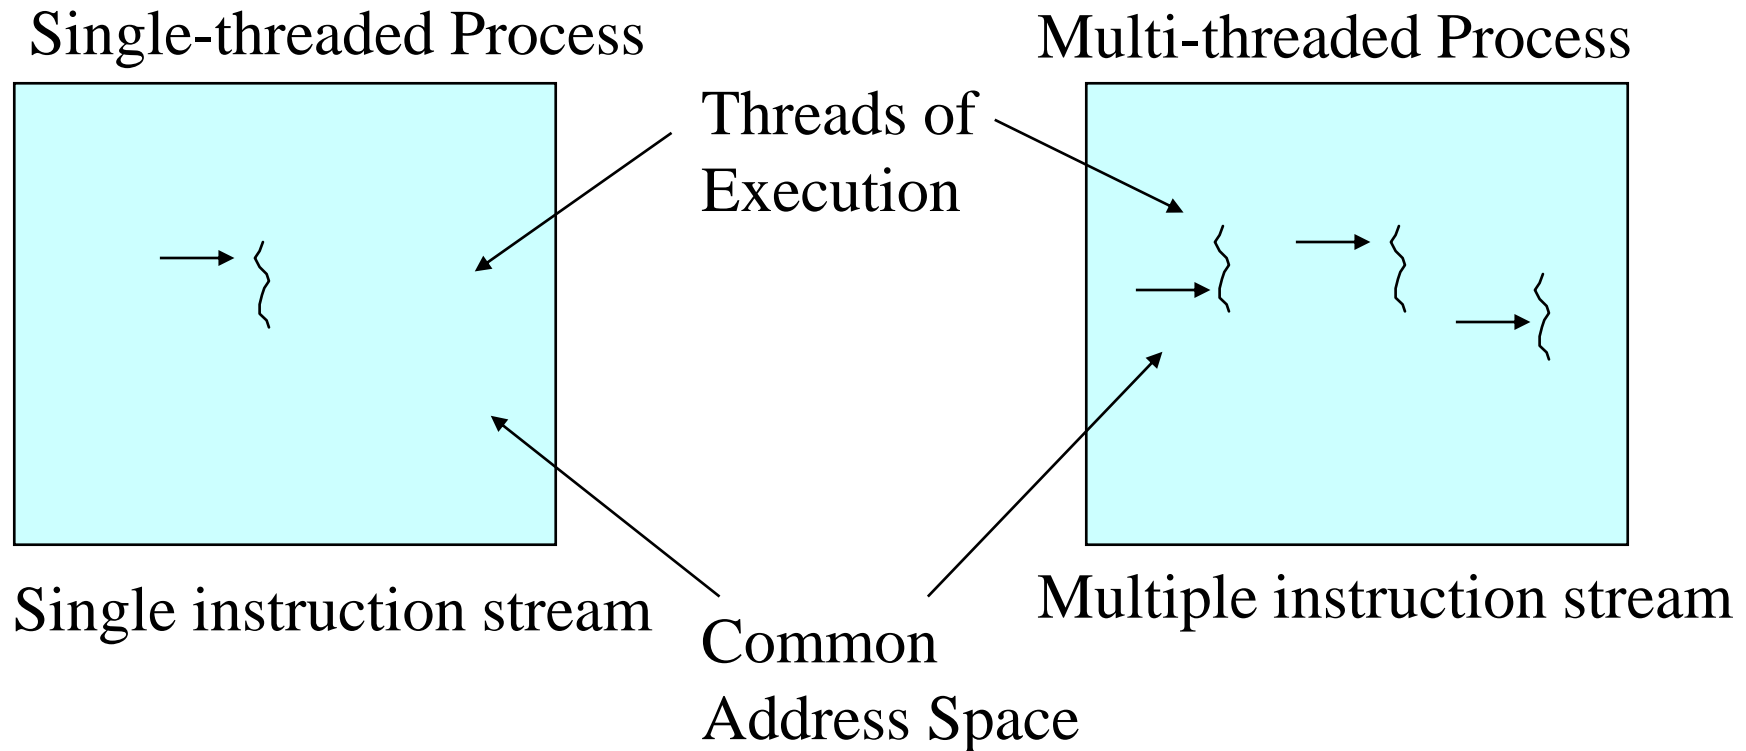
A Multithreaded Program



Threads may switch or exchange data/results

Single and Multithreaded Processes

threads are light-weight processes within a process

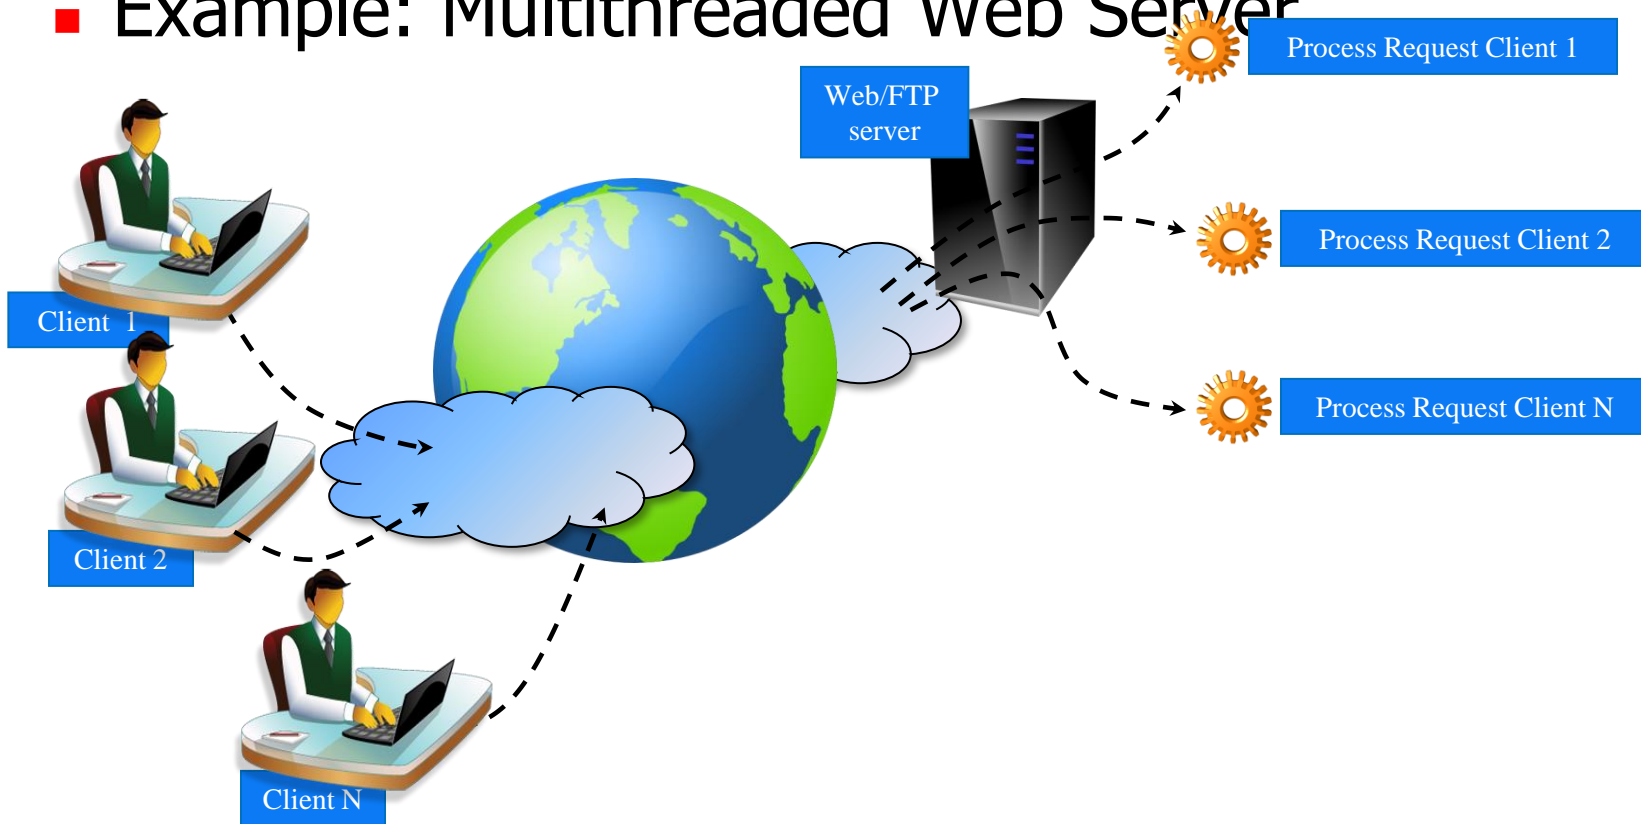


- A process?
- A thread?
- Process-based Multi-tasking vs. thread-based multitasking?

Multithreaded Server: For Serving Multiple Clients Concurrently

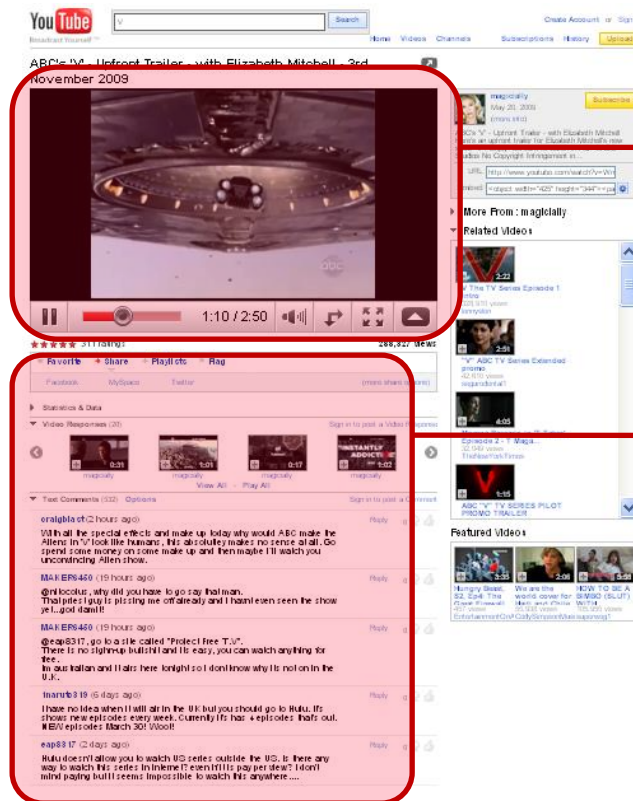
- Modern Applications

- Example: Multithreaded Web Server



Threaded Applications

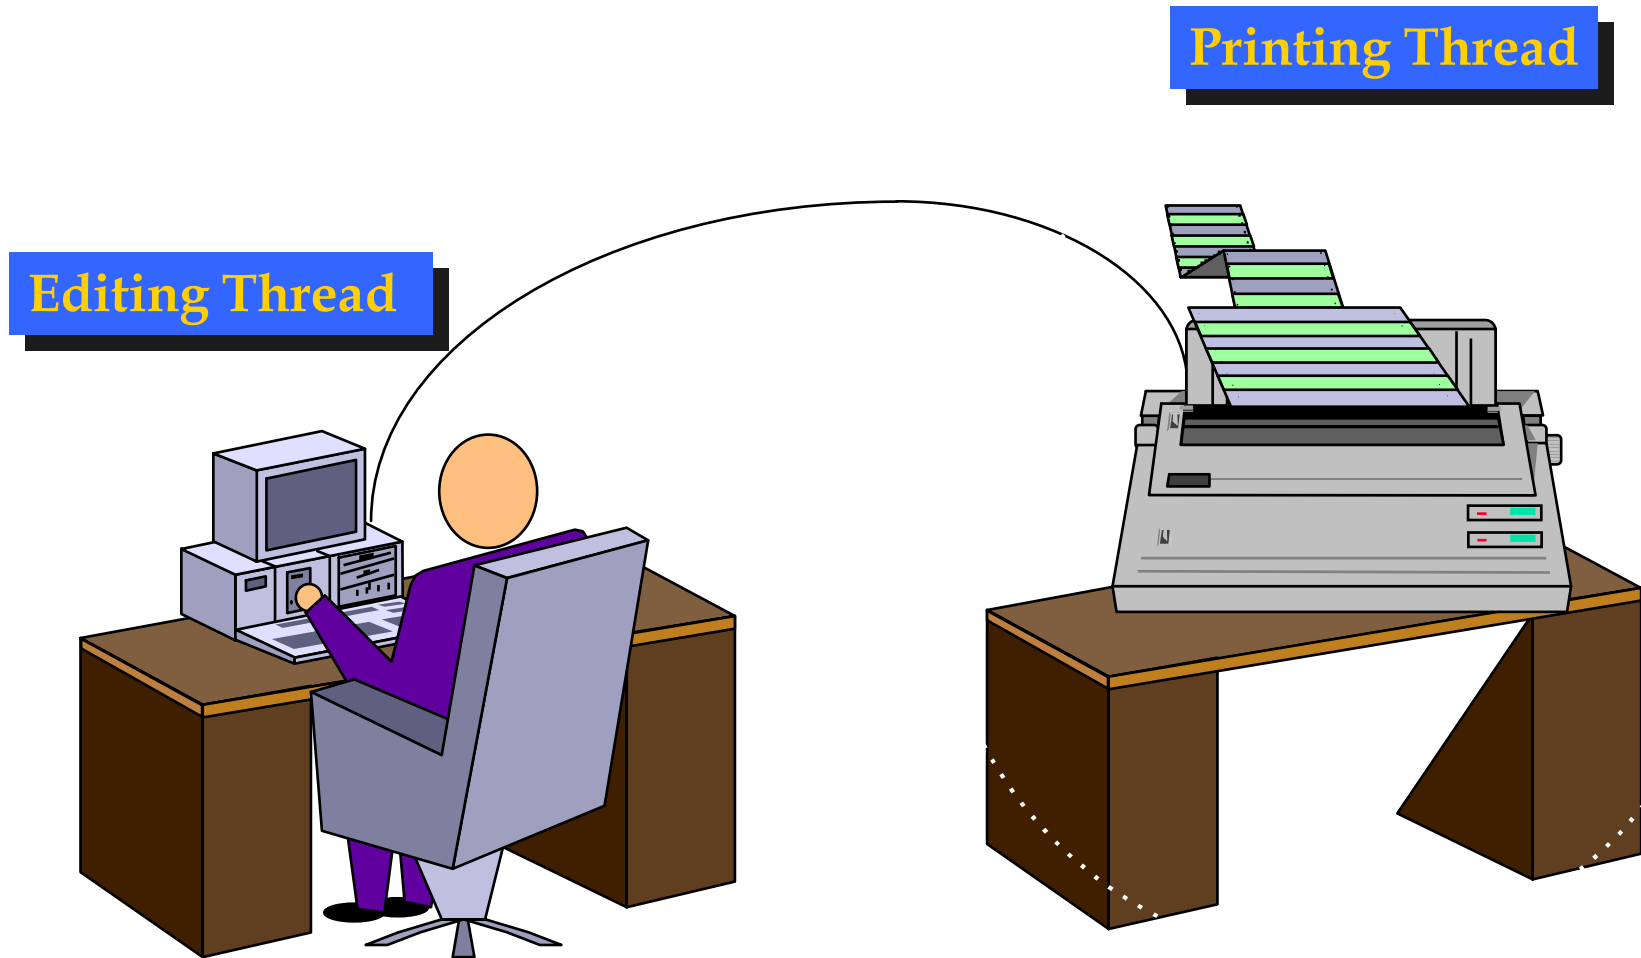
- Modern Applications
 - Example: Internet Browser + Youtube



Video Streaming

Favorites, Share,
Comments Posting

Modern Applications need Threads (ex1): Editing and Printing documents in background.



A Process

- A process consists of an execution environment together with one or more threads
- An execution environment is a collection of local kernel-managed resources to which its threads have access.
- An execution environment includes:
 - An address space.
 - Thread synchronization and communication resources (e.g., semaphores, sockets).
 - Higher-level resources such as open files and windows.

A Process


- Threads can share all resources accessible within an execution environment.
- An execution environment provides protection from threads outside it
- The data and other resources contained in an execution environment are by default inaccessible to threads residing in other execution environments.
- Example:



A Single-threaded Process

- Every software process has a single thread of execution by default.
- This is the thread that the operating system manages when it schedules the process for execution.
- This single thread has access to the program's environment and resources such as open file handles and network connections.
- As the program calls methods in objects instantiated in the code, the program's runtime stack is used to pass parameters and manage variable scopes

A Single-threaded Process



A Multi-threaded Process

- You can use programming languages to create and execute additional threads within the same process.
- Each thread is an independent sequence of execution
- What info is shared and what info is thread-specific?

A Single threaded vs. Multi-threaded Process

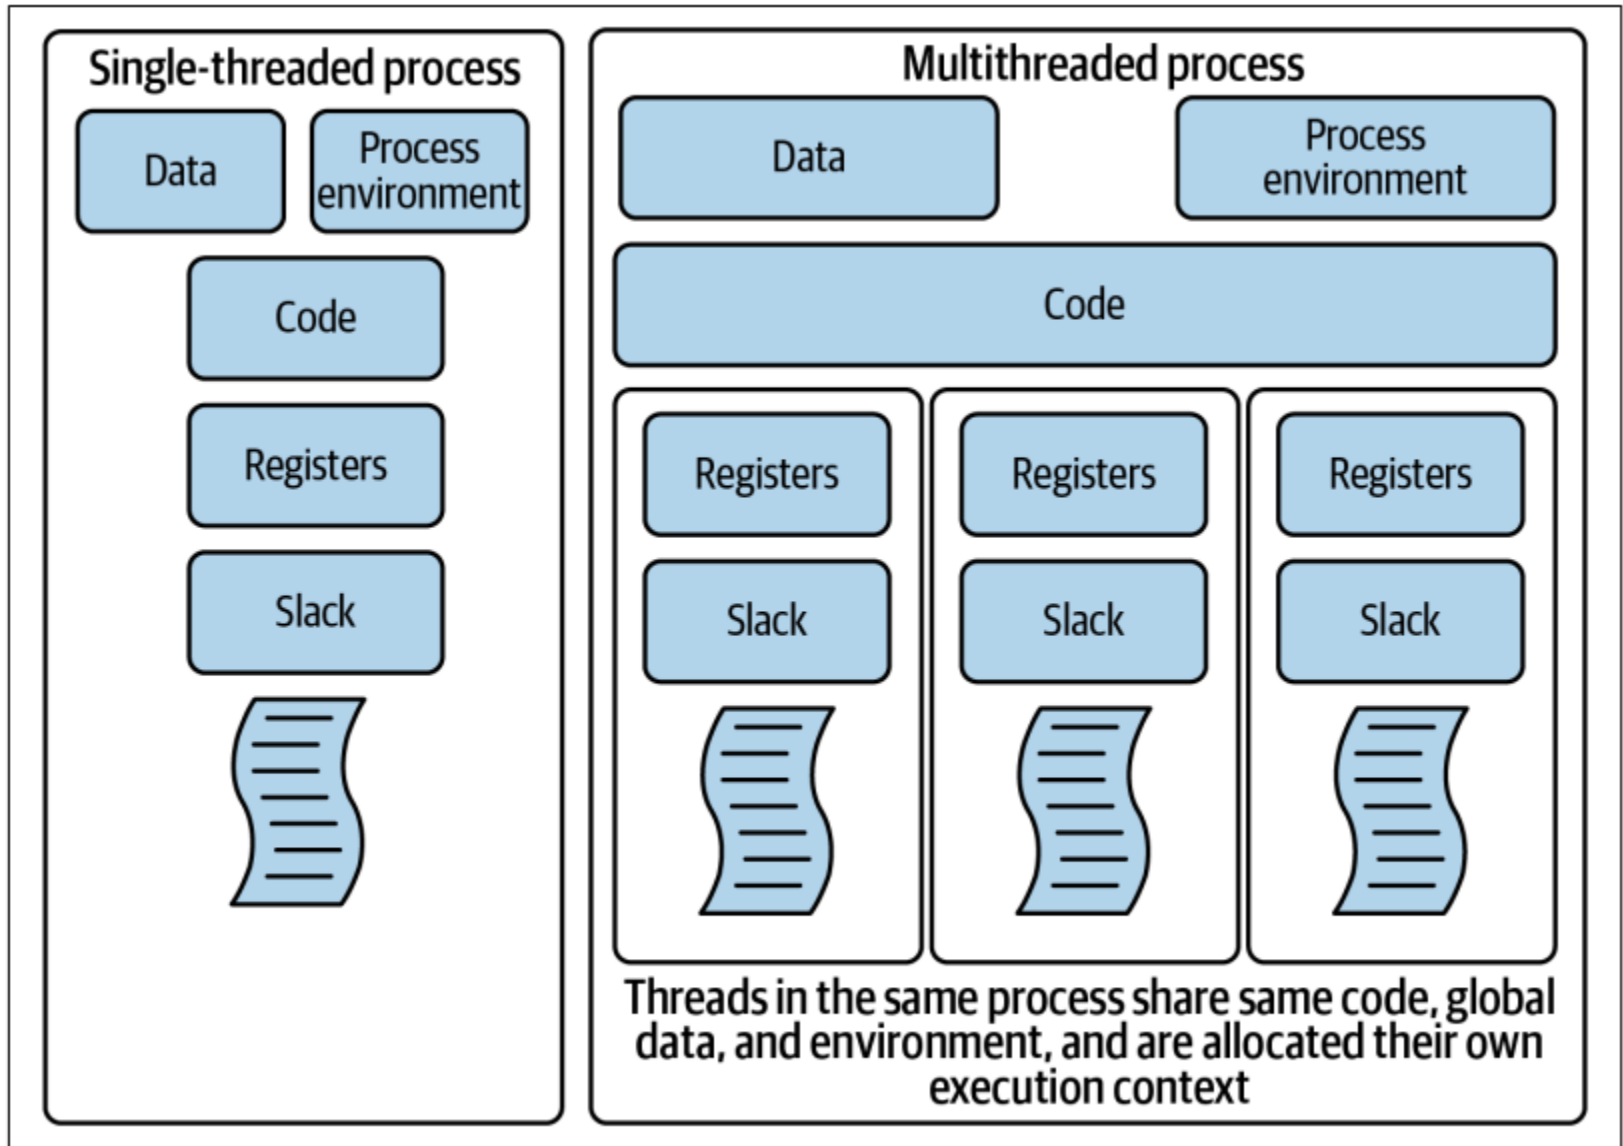


Figure 4-2. Comparing a single-threaded and multithreaded process

Process vs. Thread

- Communication
- Context switching
- Security

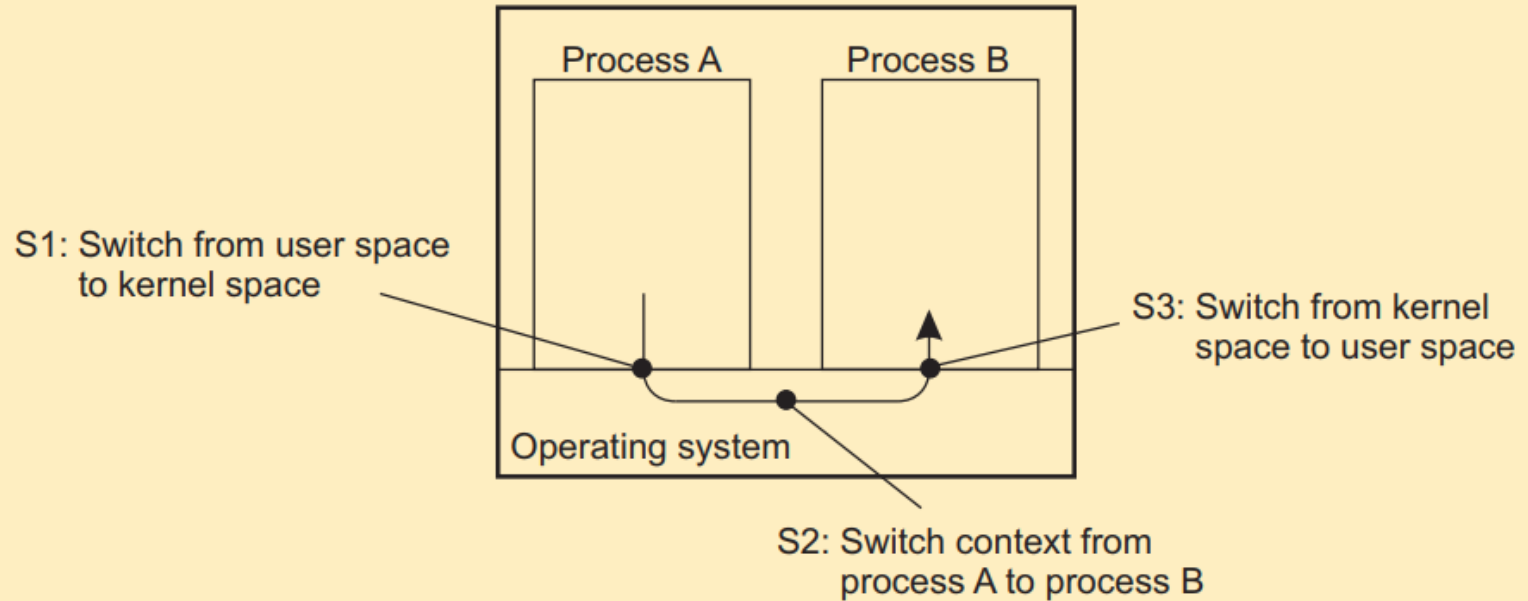
Context switching

Observations

- 1 Threads share the same address space. Thread context switching can be done entirely independent of the operating system.
- 2 Process switching is generally (somewhat) more expensive as it involves getting the OS in the loop, i.e., trapping to the kernel.
- 3 Creating and destroying threads is much cheaper than doing so for processes.

Avoid process switching

Avoid expensive context switching

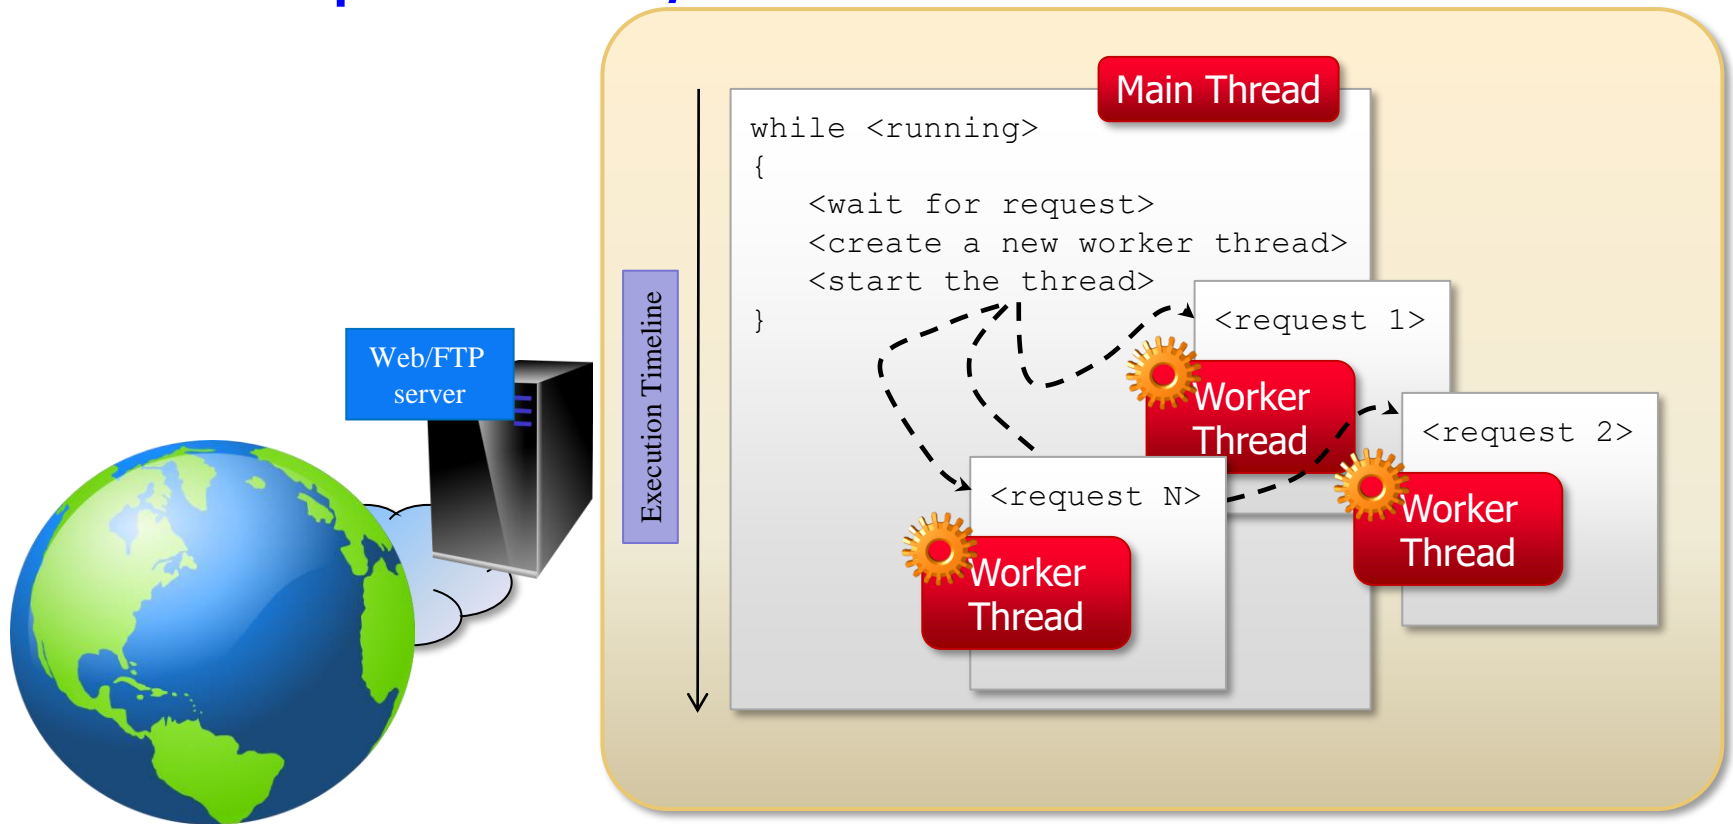


Defining Threads

- Applications – Threads are used to perform:
 - Parallelism and concurrent execution of independent tasks / operations.
 - Non blocking I/O operations.
 - Asynchronous behavior.

Defining Threads

- Example: Web/FTP Server



Defining Threads

■ Summing Up


- A Thread is a piece of code that runs in concurrent with other threads.
- Each thread is a statically ordered sequence of instructions.
- Threads are used to express concurrency on both single and multiprocessors machines.
- Programming a task having multiple threads of control – Multithreading or Multithreaded Programming.

Java Threads


- Java has built in support for Multithreading
- Synchronization
- Thread Scheduling
- Inter-Thread Communication:
 - `currentThread` `start` `setPriority`
 - `yield` `run` `getPriority`
 - `sleep` `stop` `suspend`
 - `resume`
- Java Garbage Collector is a low-priority thread.

Threading Mechanisms...

- Create a class that extends the Thread class
- Create a class that implements the Runnable interface



[a]



[b]

Which is better?

1st method: Extending Thread class

- Create a class by extending Thread class and override run() method:

```
class MyThread extends Thread
{
    public void run()
    {
        // thread body of execution
    }
}
```

- Create a thread:

```
MyThread thr1 = new MyThread();
```

- Start Execution of threads:

```
thr1.start();
```

- Create and Execute:

```
new MyThread().start();
```

An example

```
class MyThread extends Thread {  
    public void run() {  
        System.out.println(" this thread is running ... ");  
    }  
}
```

```
class ThreadEx1 {  
    public static void main(String [] args ) {  
        MyThread t = new MyThread();  
        t.start();  
    }  
}
```

2nd method: Threads by implementing Runnable interface

- Create a class that implements the interface Runnable and override run() method:

```
class MyThread implements Runnable
{
    .....
    public void run()
    {
        // thread body of execution
    }
}
```

- Creating Object:

```
MyThread myObject = new MyThread();
```

- Creating Thread Object:

```
Thread thr1 = new Thread( myObject );
```

- Start Execution:

```
thr1.start();
```

An example

```
class MyThread implements Runnable {  
    public void run() {  
        System.out.println(" this thread is running ... ");  
    }  
}
```

```
class ThreadEx2 {  
    public static void main(String [] args ) {  
        Thread t = new Thread(new MyThread());  
        t.start();  
    }  
}
```

Questions from the Previous Lecture


- Threading using Thread class versus Runnable interface?

Questions from the Previous Lecture


- Security breach within threads?
- Attackers exploit buffer overflow issues by overwriting the memory of an application
- If attackers know the memory layout of a program, they can overwrite areas that hold executable code, replacing it with their own code.
- For example, an attacker can overwrite a pointer (an object that points to another area in memory) and point it to an exploit payload, to gain control over the program

Questions from the Previous Lecture

- Security breach within threads?




(a) Process with
Single Thread



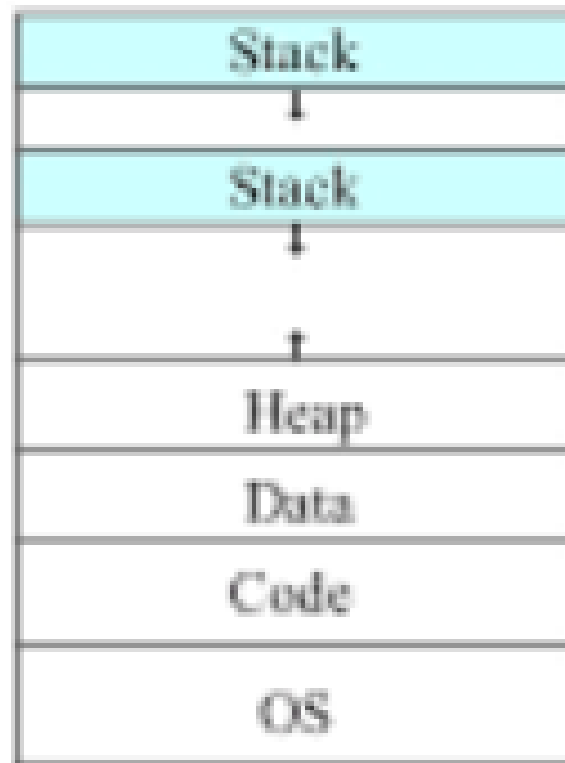
(b) Process with
Two Threads

Questions from the Previous Lecture

- Security breach within threads?



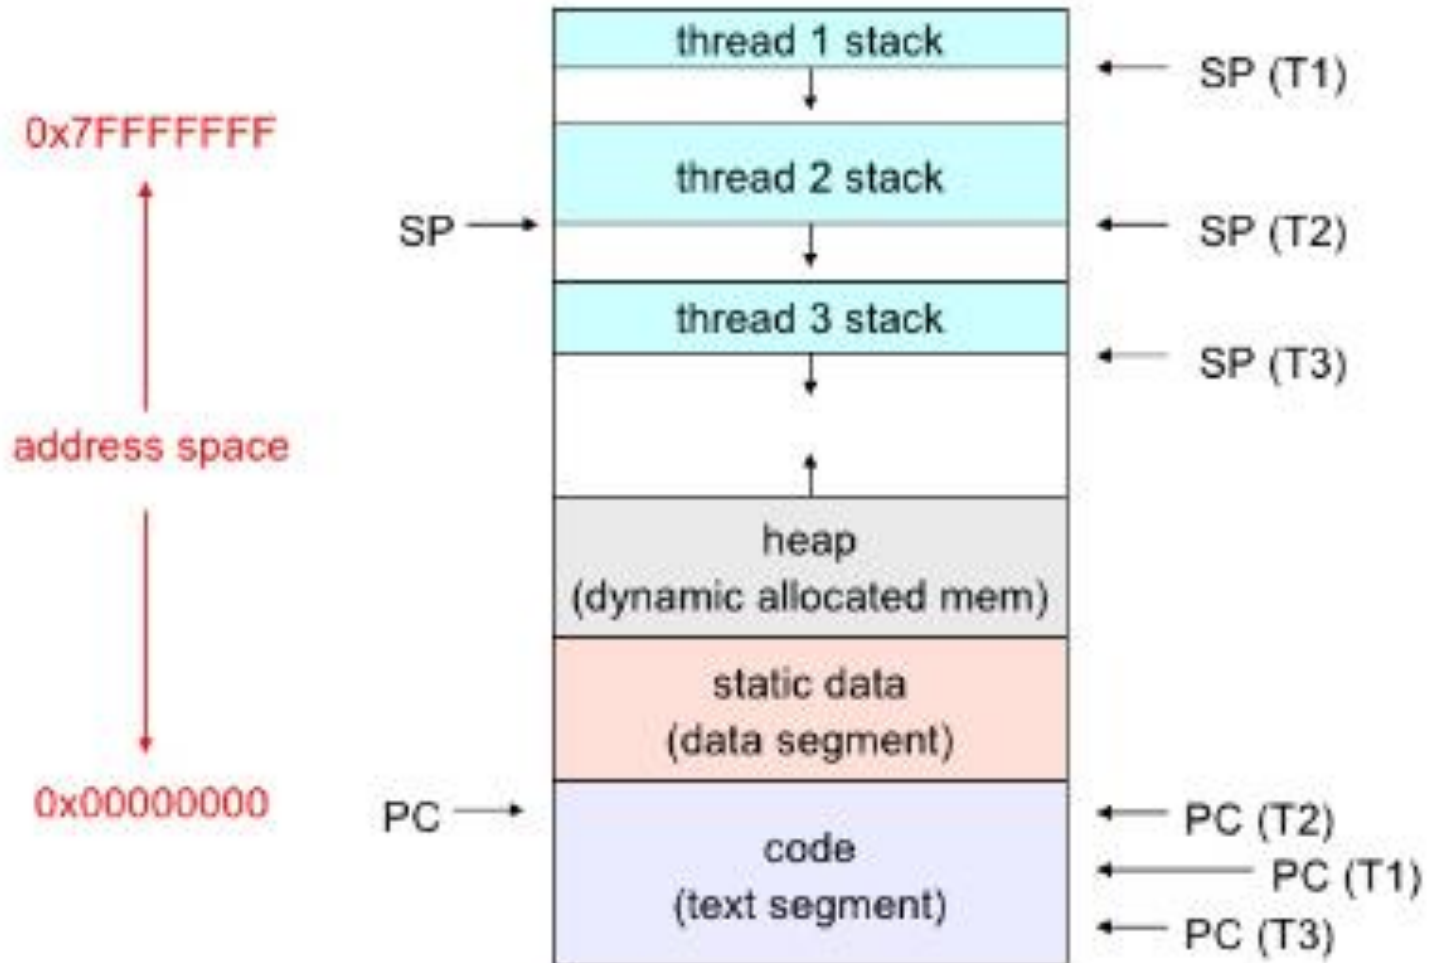
(a) Process with
Single Thread



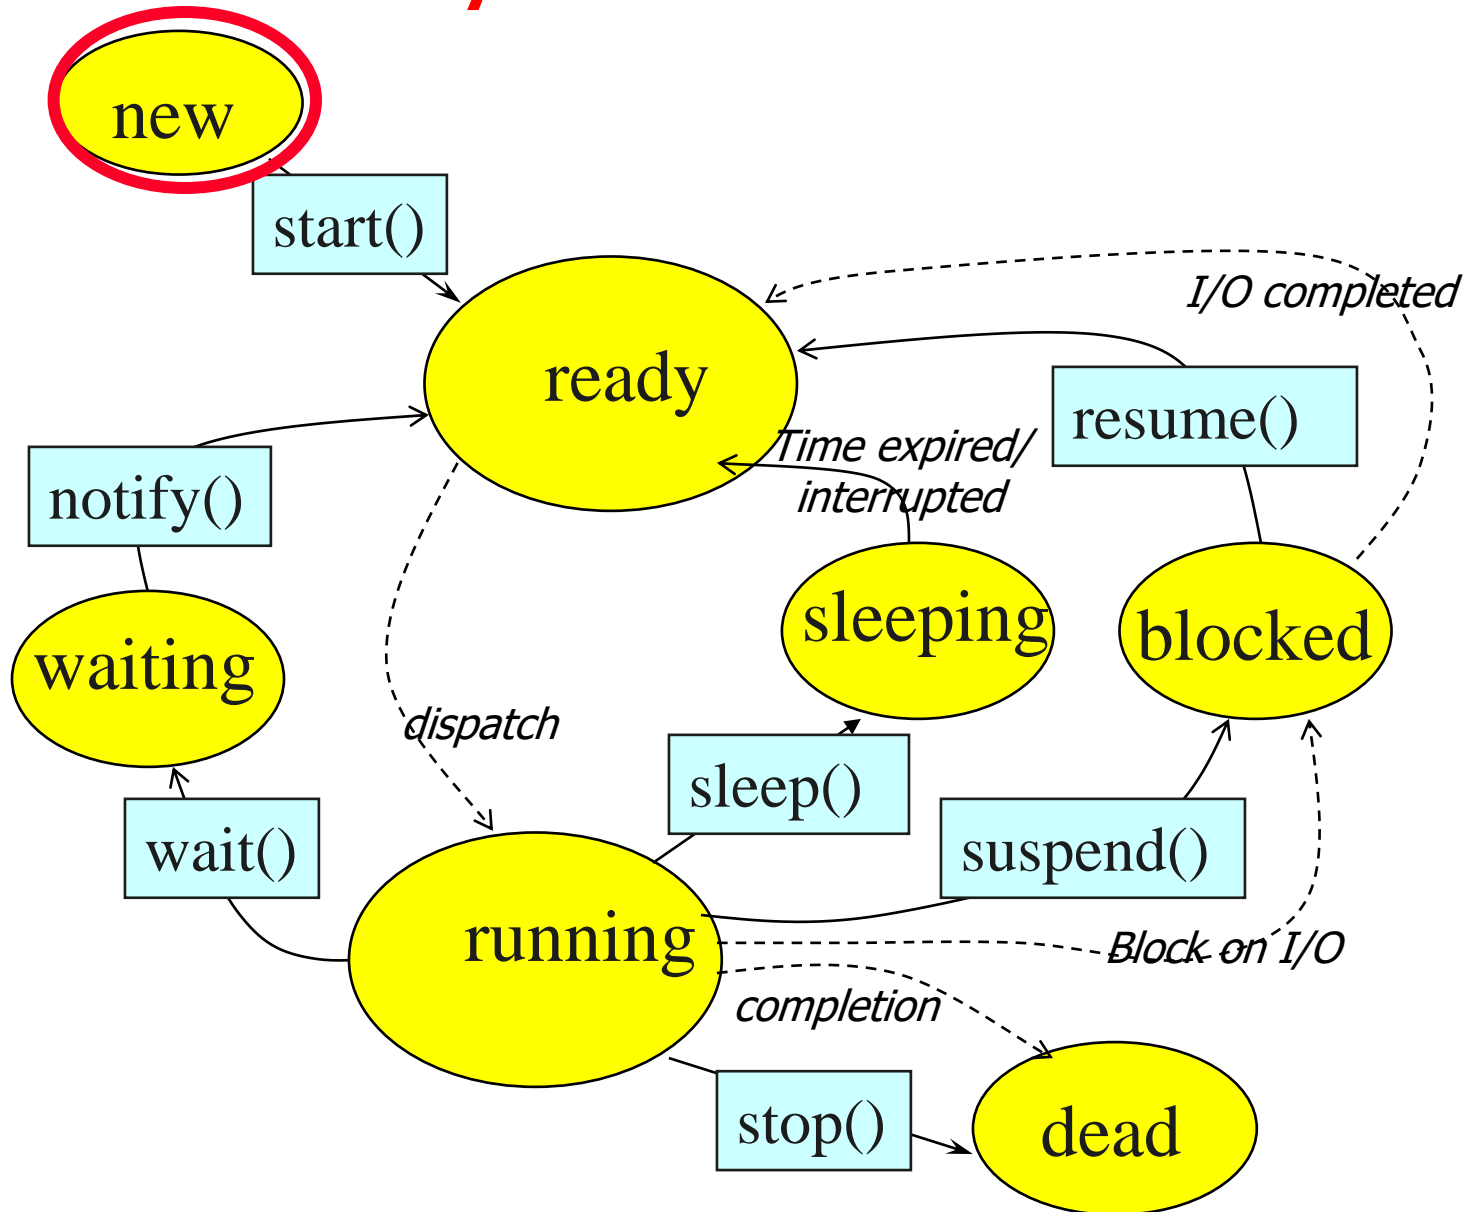
(b) Process with
Two Threads

Questions from the Previous Lecture

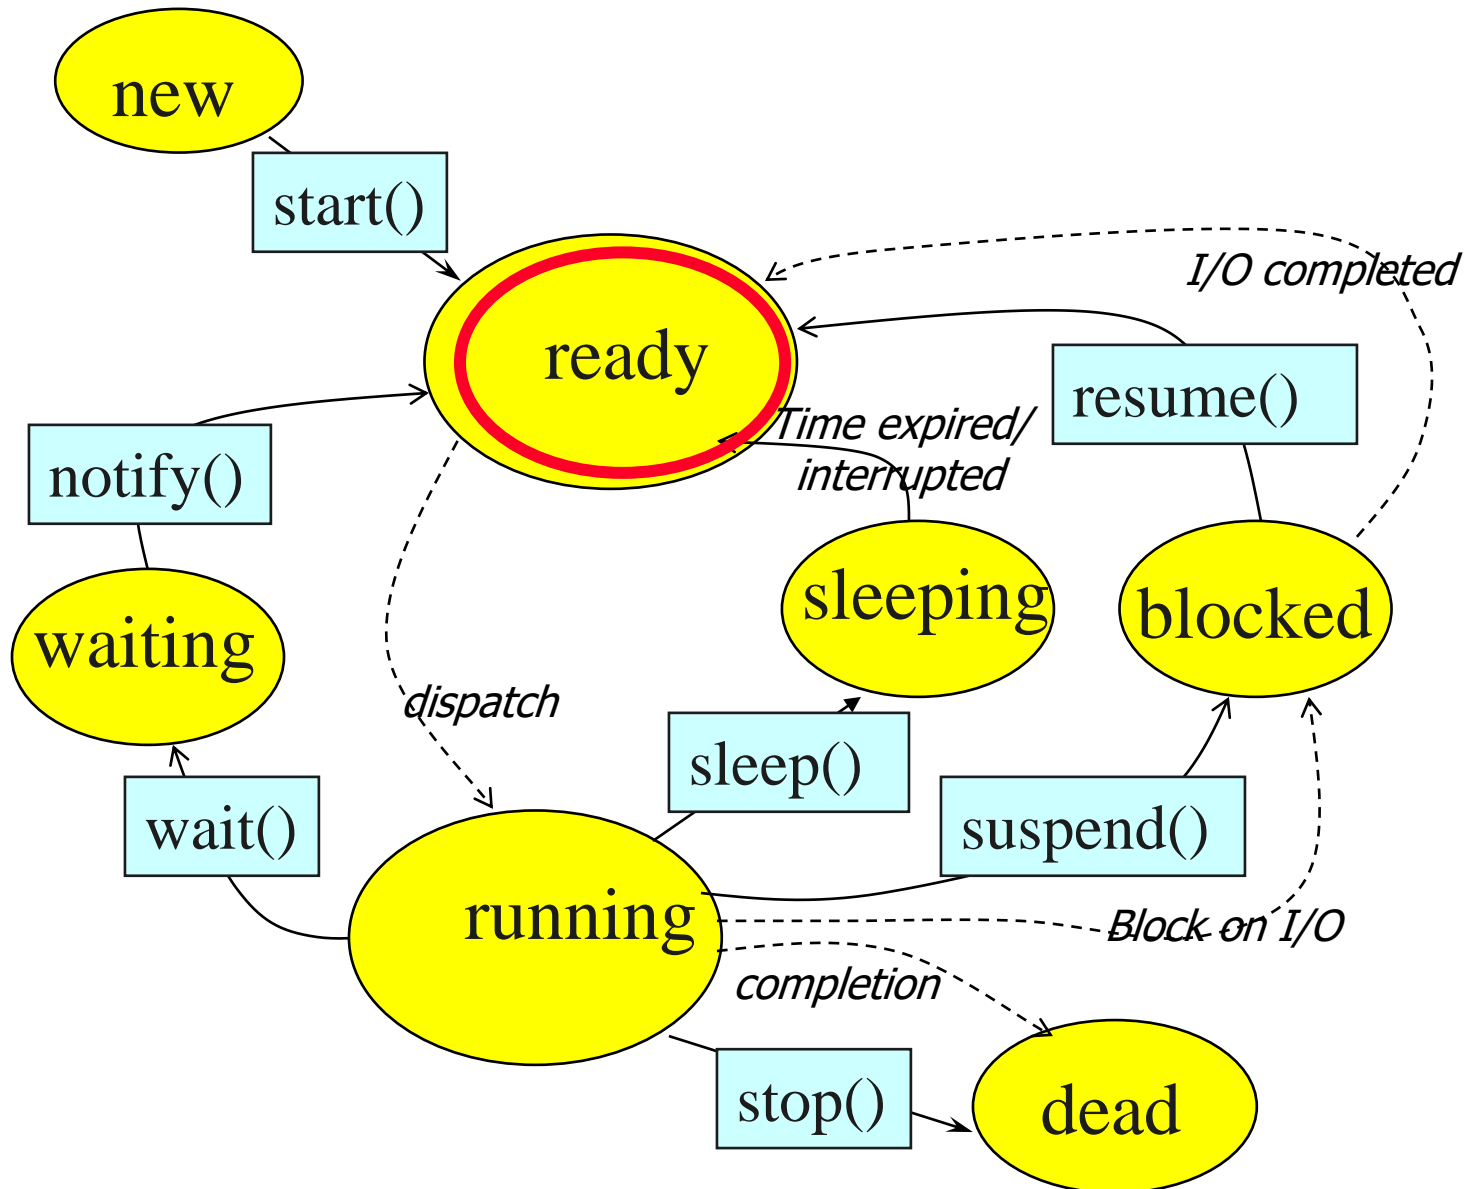
- Security breach within threads?



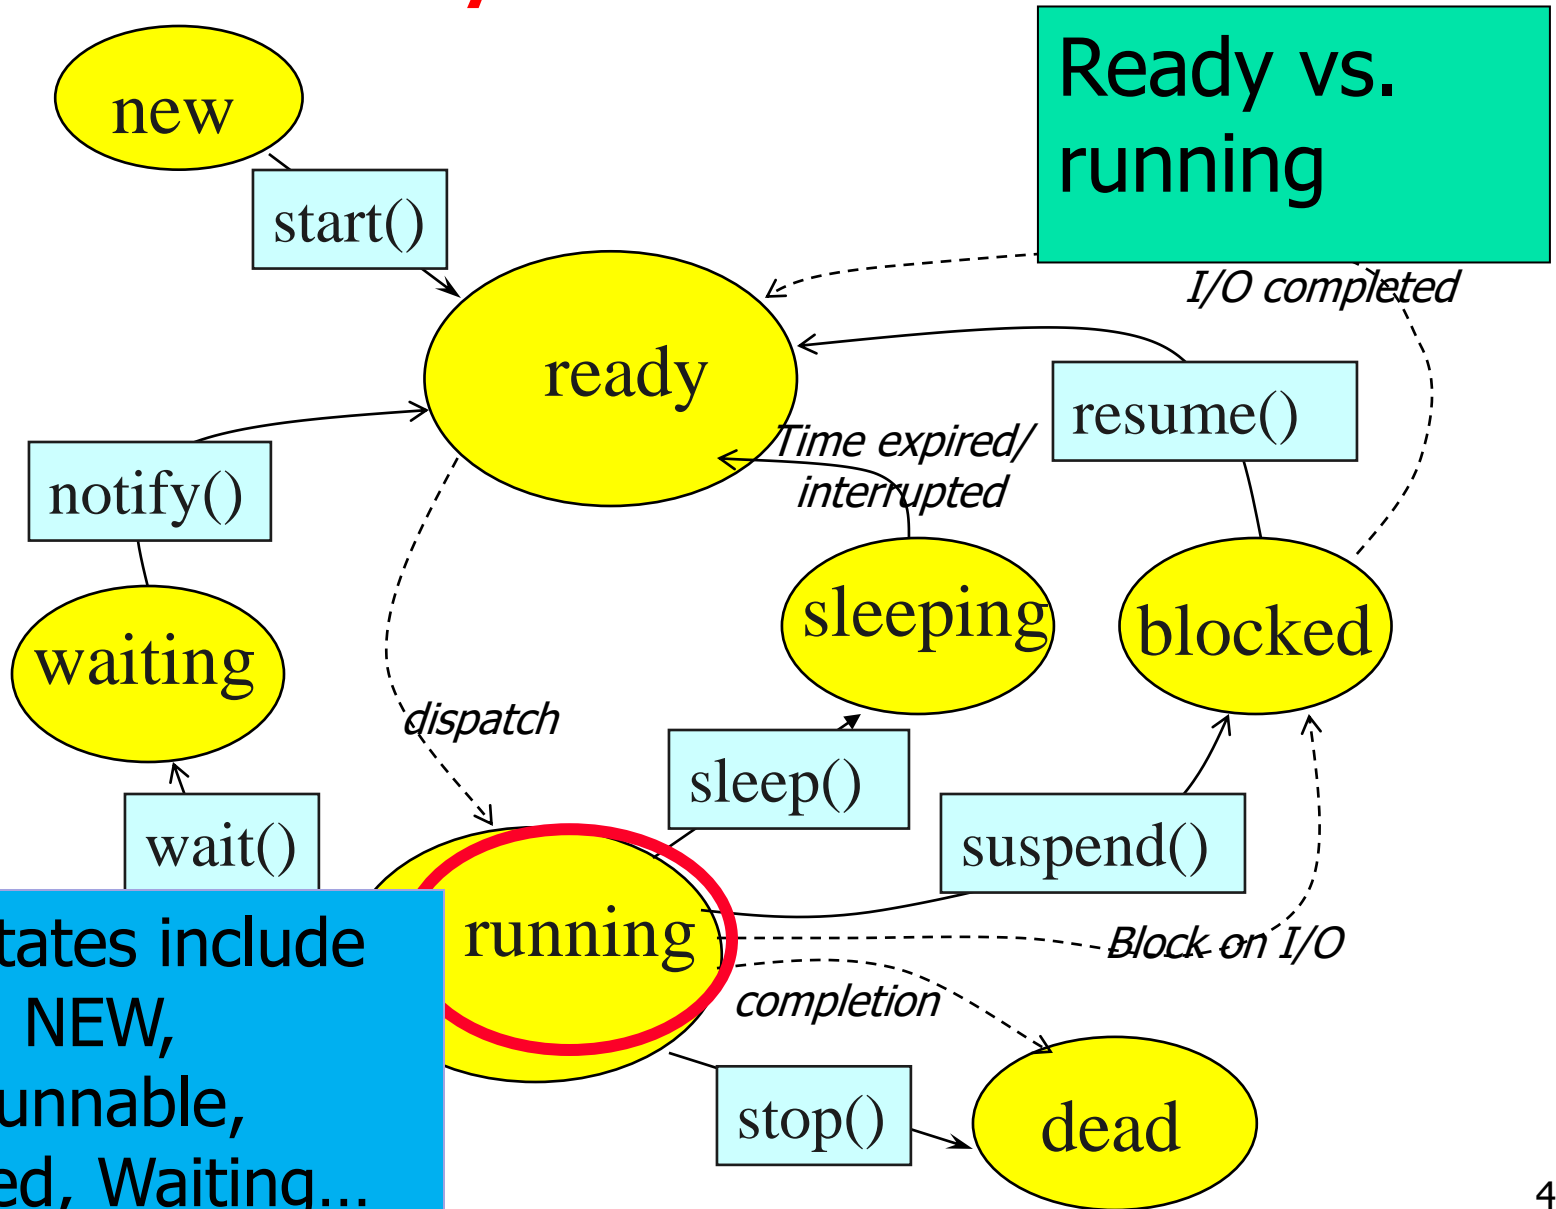
Life Cycle of Thread




Life Cycle of Thread



Life Cycle of Thread



Life Cycle of Thread



Thread Operations

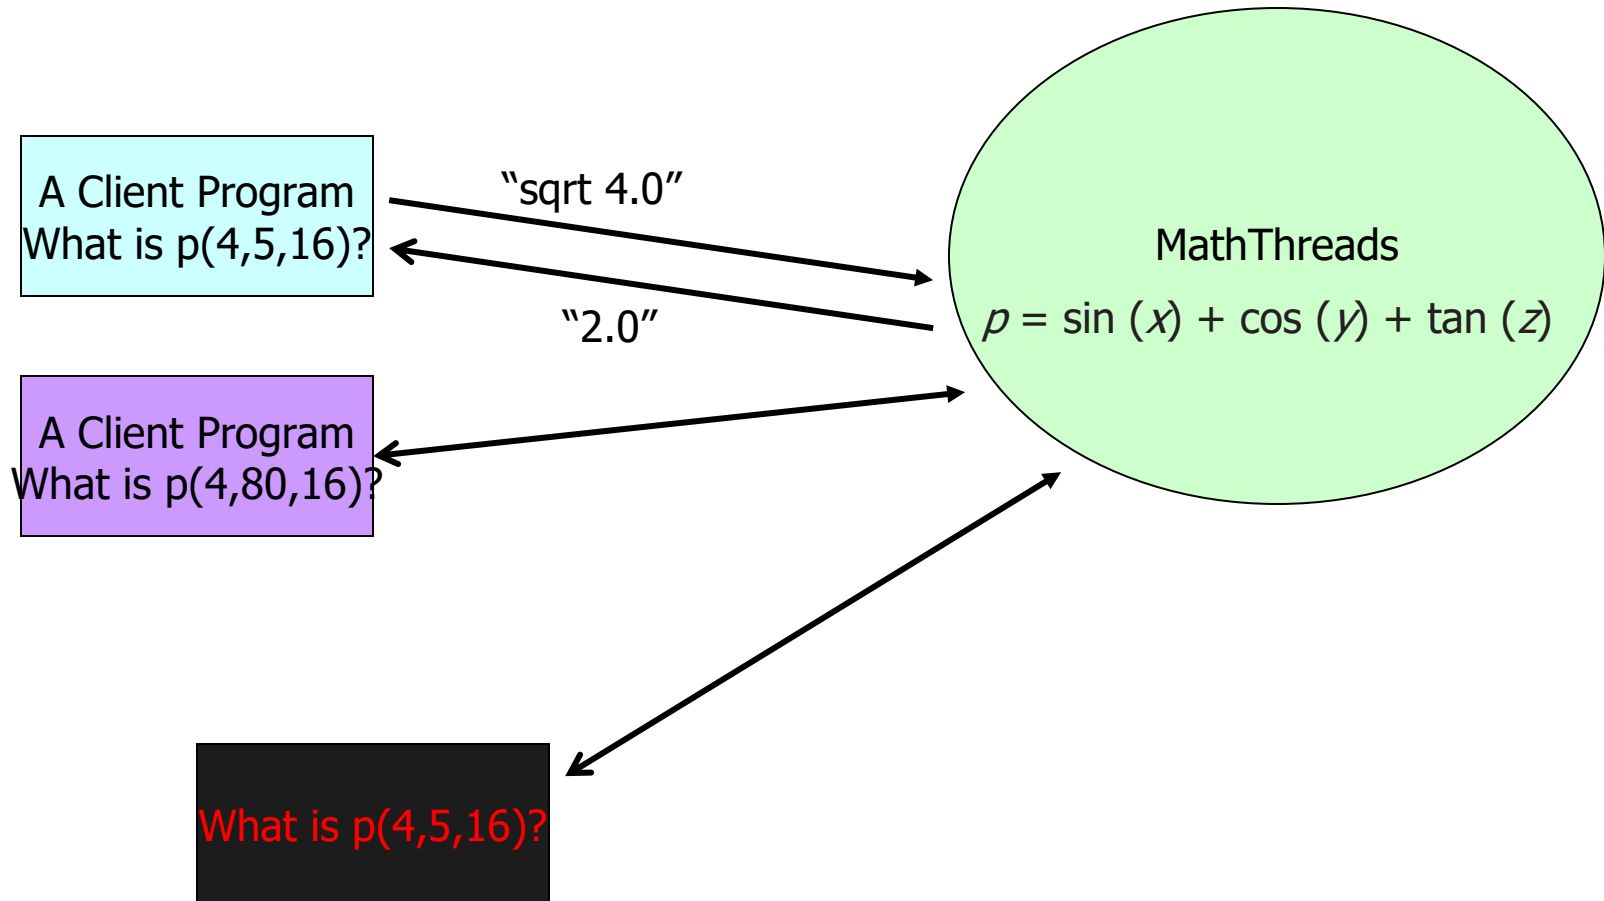
- `public static void sleep(long millis)`
throws `InterruptedException`
- For example, the code below puts the thread in sleep state for 3 minutes
- ```
try {
 Thread.sleep(3 * 60 * 1000); // thread
 sleeps for 3 minutes
} catch(InterruptedException ex){}
```


# Thread Operations

- `public static void yield()` → Back to Runnable!
- `public final boolean isAlive()`
- `void join()`  
`void join(long millis)`  
`void join(long millis, int nanos)`



# Example 1: MathServer – Demonstrates the use of Threads





```
/* MathThreads.java: A program with multiple threads performing concurrent
operations. */
import java.lang.Math;
class MathSin extends Thread {
 public double deg;
 public double res;

 public MathSin(int degree) {
 deg = degree;
 }
 public void run() {
 System.out.println("Executing sin of "+deg);
 double Deg2Rad = Math.toRadians(deg);
 res = Math.sin(Deg2Rad);
 System.out.println("Exit from MathSin. Res = "+res);
 }
}
class MathCos extends Thread {
 public double deg;
 public double res;
```

```

 public MathCos(int degree) {
 deg = degree;
 }
 public void run() {
 System.out.println("Executing cos of "+deg);
 double Deg2Rad = Math.toRadians(deg);
 res = Math.cos(Deg2Rad);
 System.out.println("Exit from MathCos. Res = "+res);
 }
}
class MathTan extends Thread {
 public double deg;
 public double res;

 public MathTan(int degree) {
 deg = degree;
 }
 public void run() {
 System.out.println("Executing tan of "+deg);
 double Deg2Rad = Math.toRadians(deg);
 res = Math.tan(Deg2Rad);
 System.out.println("Exit from MathTan. Res = "+res);
 }
}

```

```

class MathThreads {
 public static void main(String args[]) {
 MathSin st = new MathSin(45);
 MathCos ct = new MathCos(60);
 MathTan tt = new MathTan(30);
 st.start();
 ct.start();
 tt.start();
 try { // wait for completion of all thread and then sum
 st.join();
 ct.join(); //wait for completion of MathCos object
 tt.join();
 double z = st.res + ct.res + tt.res;
 System.out.println("Sum of sin, cos, tan = "+z);
 }
 catch(InterruptedException IntExp) {
 }
 }
}

```


**Run 1:**

```

[raj@mundroo] threads [1:111] java MathThreads
Executing sin of 45.0
Executing cos of 60.0
Executing tan of 30.0

```

# Multi-threaded Math Server



# Multi-threaded Math Server

```
/* MultiThreadMathServer.java: A program extending MathServer which
allows concurrent client requests and opens a new thread for each socket
connection. */
import java.net.ServerSocket;
import java.net.Socket;
public class MultiThreadMathServer
 extends MathServer implements Runnable {
 public void run() {
```

# Multi-threaded Math Server

```
 execute();
 }
 public static void main(String [] args) throws Exception {
 int port = 10000;
 if (args.length == 1) {
 try {
 port = Integer.parseInt(args[0]);
 }
 catch(Exception e) {
 }
 }
 ServerSocket serverSocket = new ServerSocket(port);
 while(true){
 //waiting for client connection
 Socket socket = serverSocket.accept();
 socket.setSoTimeout(14000);
 MultiThreadMathServer server = new MultiThreadMathServer();
 server.setMathService(new PlainMathService());
 server.setSocket(socket);
 //start a new server thread...
 new Thread(server).start();
 }
 }
}
```

How about  
the client code?



# Problems with Threads

- The basic problem in concurrent programming is coordinating the execution of multiple threads so that whatever order they are executed in, they produce the correct answer.
- Given that threads can be started and preempted **nondeterministically**, any moderately complex program will have an infinite number of possible orders of execution.
- Two main problems that concurrent programs need to avoid are **race conditions** and **deadlocks**.

# Problems with Threads (Cont'd)

## Race conditions

- Nondeterministic execution of threads implies that the code statements that comprise the threads:
  - Will execute sequentially as defined within each thread.
  - Can be overlapped in any order across threads.  
(Why?)
- When would such non-determinism cause issues?

# Problems with Threads (Cont'd)

```
public class RequestCounter {
 final static private int NUMTHREADS = 50000;
 private int count = 0;

 public void inc() {
 count++;
 }

 public int getVal() {
 return this.count;
 }

 public static void main(String[] args) throws InterruptedException {
 final RequestCounter counter = new RequestCounter();

 for (int i = 0; i < NUMTHREADS; i++) {
 // lambda runnable creation
 Runnable thread = () -> {counter.inc(); };
 new Thread(thread).start();
 }

 Thread.sleep(5000);
 System.out.println("Value should be " + NUMTHREADS + "It is: " +
 counter.getVal());
 }
}
```

# Problems with Threads (Cont'd)


## Race conditions

- To perform an increment of a counter, the CPU must:
  - Load the current value into a register.
  - Increment the register value.
  - Write the results back to the original memory location
- This is a sequence of **three** machine level instructions.
- However, such operations are not treated as a **single atomic operation**.

# Problems with Threads (Cont'd)

## Race conditions

| Thread 1                     | Thread 2                     | Thread 1                     | Thread 2                     |
|------------------------------|------------------------------|------------------------------|------------------------------|
| Reads (x) into register      |                              | Reads (x) into register      |                              |
| Register value +6            |                              | Register value +6            |                              |
| Writes register value to (x) |                              |                              | Reads (x) into register      |
|                              | Reads (x) into register      |                              | Register value +1            |
|                              | Register value +1            |                              | Writes register value to (x) |
|                              | Writes register value to (x) | Writes register value to (x) |                              |






Figure 4-3. Increments are not atomic at the machine level

# Problems with Threads (Cont'd)

## Deadlocks

- To ensure correct results in multithreaded code, we have to restrict the inherent nondeterminism to serialize access to critical sections.
- This avoids race conditions.
- But, we can write code that restricts nondeterminism so much that our program stops. And never continues. This is formally known as a deadlock..
- How would it happen within threads?

# Required Readings

- CDK Book (Text Book)
  - Chapter 7 – “Operating System Support”
- Chapter 4: An Overview of Concurrent Systems, from: Gorton, Ian. *Foundations of Scalable Systems*. " O'Reilly Media, Inc.", 2022.
- Chapter 14: Multithread Programming, from R. Buyya, S. Selvi, X. Chu, “**Object Oriented Programming with Java: Essentials and Applications**”, McGraw Hill, New Delhi, India, 2009.