

Machine learning

Presented by : Dr. Hanaa Bayomi



Lecture 2 : Linear Regression

LINEAR REGRESSION WITH ONE VARIABLE


- Model Representation
- Cost Function
- Gradient Descent

MODEL REPRESENTATION

Housing Prices (Portland, OR)

dependent
variable

Price
(in 1000s
of dollars)



Independent
variable

Size (feet²)

Supervised Learning

Regression:

“right answers” or “Labeled data” given

Predict continuous valued output (price)

MODEL REPRESENTATION

Training set of
housing prices
(Portland, OR)

Size in feet² (x)

Price (\$) in 1000's (y)

2104

460

1416

232

1534

315

852

178

...

...

m

Notation:

m = Number of training examples

x's = "input" variable / features

y's = "output" variable / "target" variable

(**x**,**y**) one training example (one row)

(**x**⁽ⁱ⁾,**y**⁽ⁱ⁾) **i**th training example

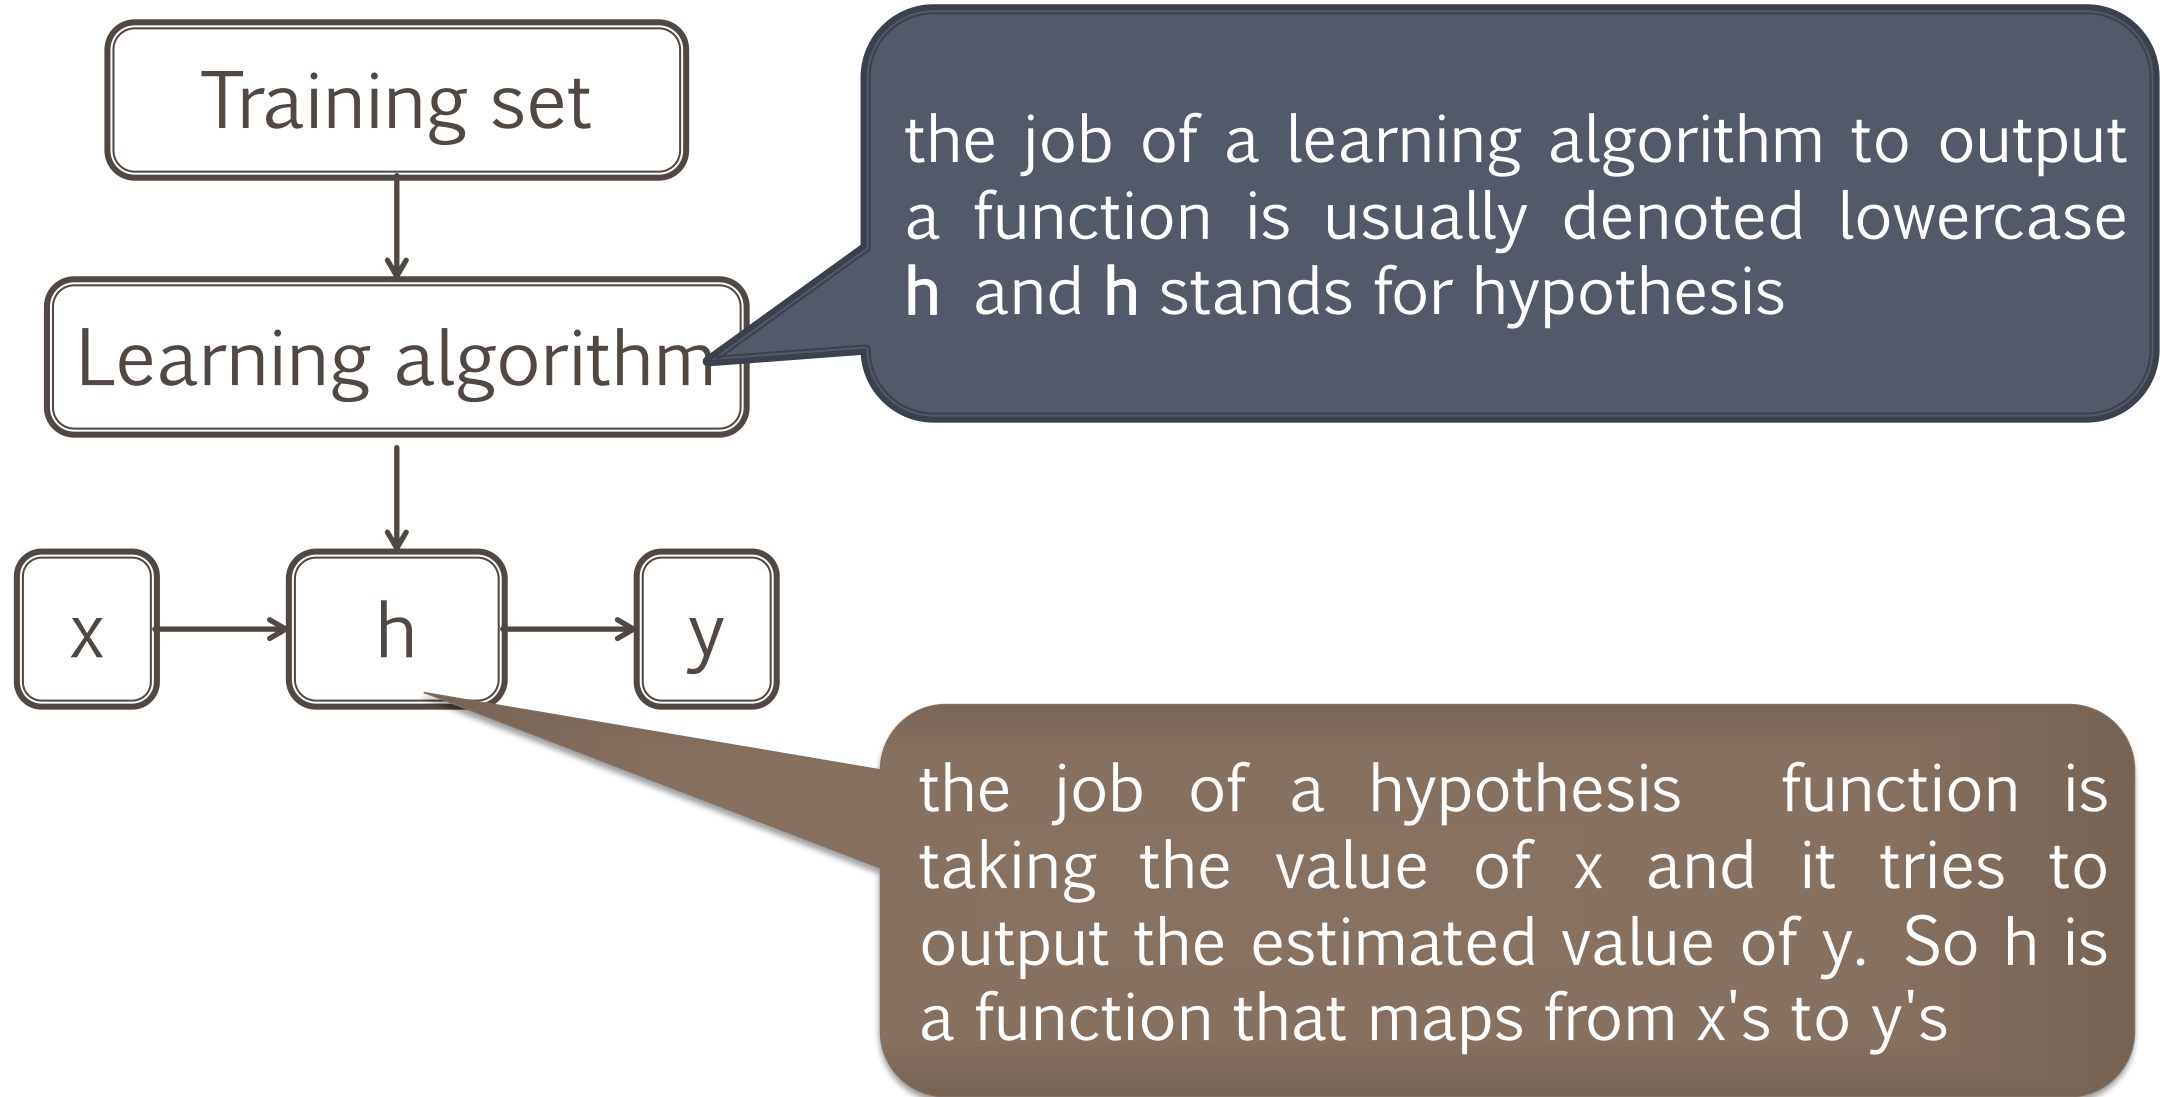
Example

x⁽¹⁾ 2104

y⁽²⁾ 232

x⁽⁴⁾ 852


MODEL REPRESENTATION




MODEL REPRESENTATION

How do we represent h ?

$$h(x) = \theta_0 + \theta_1 x$$




Linear Equations




Types of Regression Models


Positive Linear Relationship




Relationship NOT Linear



Negative Linear Relationship



No Relationship



COST FUNCTION

- *The cost function*, let us figure out how to fit the best possible straight line to our data.


Training Set	Size in feet ² (x)	Price (\$) in 1000's (y)
	2104	460
	1416	232
	1534	315
	852	178

Hypothesis: $h_{\theta}(x) = \theta_0 + \theta_1 x$

How to choose θ_i 's ?


Scatter plot

- 1. Plot of All (X_i, Y_i) Pairs
- 2. Suggests How Well Model Will Fit




Thinking Challenge

How would you draw a line through the points?
How do you determine which line 'fits best'?



Thinking Challenge

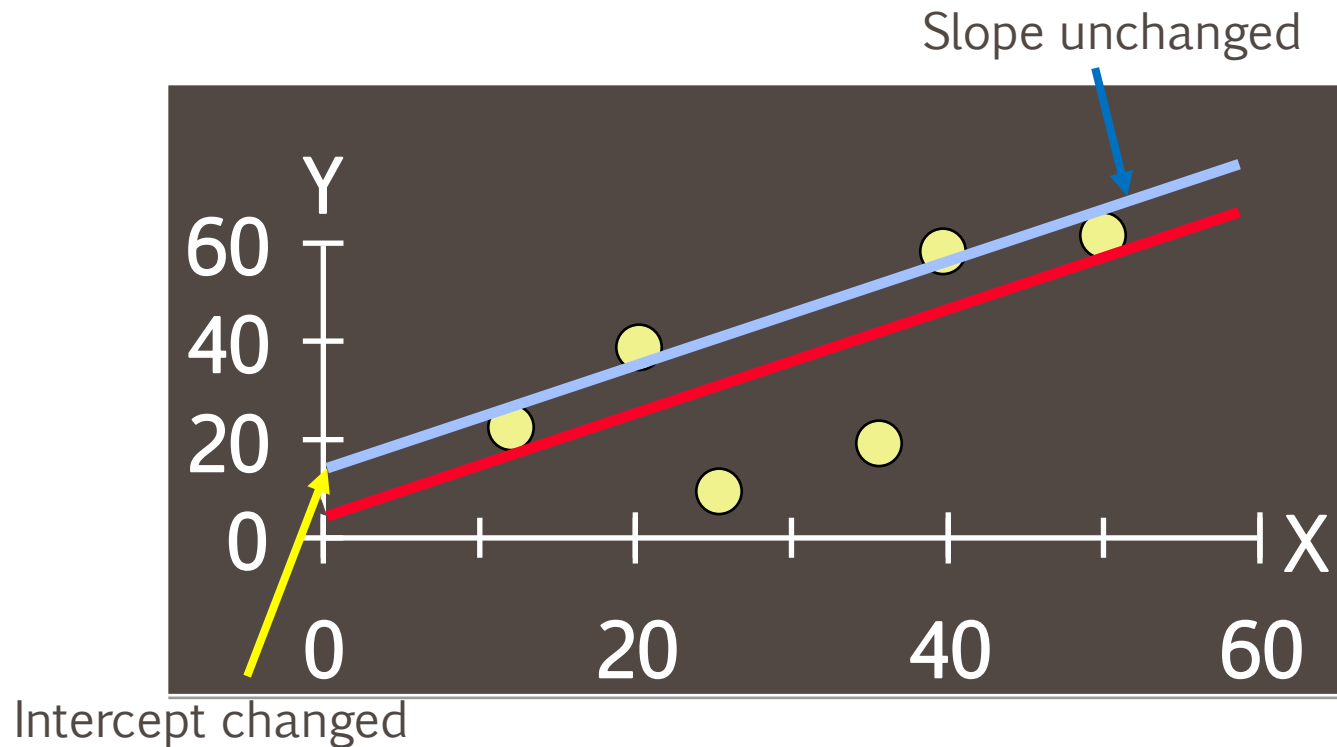
How would you draw a line through the points?
How do you determine which line 'fits best'?



Thinking Challenge


How would you draw a line through the points?

How do you determine which line 'fits best'?



Thinking Challenge

How would you draw a line through the points?
How do you determine which line 'fits best'?



Least Squares

- 1. 'Best Fit' Means Difference Between Actual Y Values and Predicted Y Values is a Minimum. So square errors!

$$\sum_{i=1}^m (Y_i - h\theta(x_i))^2 = \sum_{i=1}^m \hat{\varepsilon}_i^2$$

Least Squares


- 1. 'Best Fit' Means Difference Between Actual Y Values & Predicted Y Values Are a Minimum. So square errors!

$$\sum_{i=1}^m (Y_i - h\theta(x_i))^2 = \sum_{i=1}^m \hat{\epsilon}_i^2$$



- 2. LS Minimizes the Sum of the Squared Differences (errors) (SSE)

Least Squares Graphically

LS minimizes $\sum_{i=1}^n \hat{\varepsilon}_i^2 = \hat{\varepsilon}_1^2 + \hat{\varepsilon}_2^2 + \hat{\varepsilon}_3^2 + \hat{\varepsilon}_4^2$




Least Squared **errors** Linear Regression



$$\text{minimize}_{\theta_0, \theta_1} \frac{1}{2m} \sum_{i=1}^m \left(h_{\theta}(x^{(i)}) - y^{(i)} \right)^2$$

COST FUNCTION



Idea: Choose θ_0, θ_1 so that $h_\theta(x)$ is close to y for our training examples (x, y)

$$\text{Minimize}_{\theta_0 \theta_1} \frac{1}{2m} \sum_i^m (h_\theta(x^i) - y^i)^2$$



$$h_\theta(x^i) = \theta_0 + \theta_1 x^i$$

$h_\theta(x^i)$ predictions on the training set

y^i the actual values

$$j(\theta_0, \theta_1) = \frac{1}{2m} \sum_i^m (h_\theta(x^i) - y^i)^2$$

$$\text{Minimize}_{\theta_0 \theta_1} j(\theta_0, \theta_1)$$

Cost function visualization

Consider a simple case of hypothesis by setting $\theta_0=0$, then \mathbf{h} becomes :
 $\mathbf{h}_\theta(\mathbf{x})=\theta_1\mathbf{x}$

Each value of θ_1 corresponds to a different hypothesis as it is the **slope** of the line


which corresponds to different lines passing through the **origin** as shown in plots below as **y-intercept** i.e. θ_0 is nulled out.

$$J(\theta_1) = \frac{1}{2m} \sum_{i=1}^m \left(\theta_1 x^{(i)} - y^{(i)} \right)^2$$

$$\text{At } \theta_1=2, \quad J(2) = \frac{1}{2 * 3} (1^2 + 2^2 + 3^2) = \frac{14}{6} = 2.33$$


$$\text{At } \theta_1=1, \quad J(1) = \frac{1}{2 * 3} (0^2 + 0^2 + 0^2) = 0$$

$$\text{At } \theta_1=0.5, \quad J(0.5) = \frac{1}{2 * 3} (0.5^2 + 1^2 + 1.5^2) = 0.58$$



Cost function visualization

$$J(\theta_1) = \frac{1}{2m} \sum_{i=1}^m (\theta_1 x^{(i)} - y^{(i)})^2$$




At $\theta_1=2$, $J(2) = \frac{1}{2 * 3} (1^2 + 2^2 + 3^2) = \frac{14}{6} = 2.33$

At $\theta_1=1$, $J(1) = \frac{1}{2 * 3} (0^2 + 0^2 + 0^2) = 0$

At $\theta_1=0.5$, $J(0.5) = \frac{1}{2 * 3} (0.5^2 + 1^2 + 1.5^2) = 0.58$

On **plotting points** like this further, one gets the following graph for the cost function which is dependent on parameter θ_1 .

plot each value of θ_1 corresponds to a different hypothesizes



Cost function visualization

$$J(\theta_1) = \frac{1}{2m} \sum_{i=1}^m (\theta_1 x^{(i)} - y^{(i)})^2$$

What is the optimal value of θ_1 that minimizes $J(\theta_1)$?


It is clear that best value for $\theta_1 = 1$ as $J(\theta_1) = 0$, which is the minimum.

How to find the best value for θ_1 ?

Plotting ?? Not practical specially in high dimensions?


The solution :

1. Analytical solution: not applicable for large datasets
2. Numerical solution: ex: Gradient descent .



plotting the cost function $j(\theta_0, \theta_1)$

- Previously we plotted our cost function by plotting
 - θ_1 vs $J(\theta_1)$
- Now we have two parameters
 - Plot becomes a bit more complicated
 - Generates a 3D surface plot where axis are
 - $X = \theta_1$
 - $Z = \theta_0$
 - $Y = J(\theta_0, \theta_1)$



COST FUNCTION (RECAP)

Hypothesis:

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

Parameters:

$$\theta_0, \theta_1$$

Cost Function:

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Goal: minimize $J(\theta_0, \theta_1)$
 θ_0, θ_1

Gradient Descent

GRADIENT DESCENT

- Iterative solution not only in linear regression. It's actually used all over the place in machine learning.
- Objective: minimize any function (Cost Function J)

PROBLEM SETUP


Have some function $J(\theta_0, \theta_1)$


Want $\min_{\theta_0, \theta_1} J(\theta_0, \theta_1)$

Outline:

- Start with some θ_0, θ_1
- Keep changing θ_0, θ_1 to reduce $J(\theta_0, \theta_1)$
until we hopefully end up at a minimum

Imagine that this is a landscape of grassy park, and you want to go to the lowest point in the park as rapidly as possible





With different starting point

Gradient descent Algorithm

repeat until convergence $\{\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1) \forall j \in \{0, 1\}\}$

- Where

- $:=$ is the assignment operator
- α is the **learning rate** which basically defines how big the steps are during the descent
- $\frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$ is the **partial derivative** term
- $j = 0, 1$ represents the **feature index number**


Also the parameters should be **updated simulatenously**, i.e. ,

$$temp_0 := \theta_0 - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1)$$


$$temp_1 := \theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1)$$

$$\theta_0 := temp_0$$

$$\theta_1 := temp_1$$




EXAMPLE




$$\theta_1 = \theta_1 - \alpha \frac{d}{d\theta_1} j(\theta_1)$$

$$\theta_1 = \theta_1 - \alpha (+ve) \downarrow$$



$$\theta_1 = \theta_1 - \alpha (-ve) \uparrow$$

QUESTION



Current value of θ_1

$$\theta_1 := \theta_1 - \alpha \frac{d}{d\theta_1} J(\theta_1)$$

GRADIENT DESCENT FOR A LINEAR REGRESSION

Gradient descent algorithm

repeat until convergence {
 $\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$
 (for $j = 1$ and $j = 0$)
}

Linear Regression Model

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

$$\frac{d}{d\theta_j} j(\theta_0, \theta_1) = \frac{d}{d\theta_j} \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x_i) - Y_i)^2$$

$$\frac{d}{d\theta_j} j(\theta_0, \theta_1) = \frac{d}{d\theta_j} \frac{1}{2m} \sum_{i=1}^m (\theta_0 + \theta_1(x_i) - Y_i)^2$$

$$j = 0 : \frac{d}{d\theta_0} j(\theta_0, \theta_1) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x_i) - Y_i)$$

$$j = 1 : \frac{d}{d\theta_1} j(\theta_0, \theta_1) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x_i) - Y_i) \bullet x_i$$

G.D. FOR LINEAR REGRESSION

Gradient descent algorithm

repeat until convergence {

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})$$

$$\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x^{(i)}$$

}

“Batch” Gradient Descent

“Batch”: Each step of gradient descent uses all the training examples.

repeat until convergence {

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})$$

$$\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x^{(i)}$$

}

Linear Regression

Using

TensorFlow

1-D Data Example

Data Preparation

```
import numpy as np
```


```
num_of_points = 100          #Generate 100 Data Points  
points = []  
for i in range(num_of_points):  
    x1= np.random.normal(0.0, 0.55)  
    y1= x1 * 0.1 + 0.3 + np.random.normal(0.0, 0.01)  
    points.append([x1, y1])  
x_data = [v[0] for v in points]  
y_data = [v[1] for v in points]
```


Draw Data

```
import matplotlib.pyplot as plt
```

```
plt.plot(x_data, y_data, 'ro', label='Original data')  
plt.legend()  
plt.show()
```

Original Data



Variables and Nodes Preparation

```
import tensorflow as tf

#initialize weights "W and bias "b"
W = tf.Variable(tf.random_uniform([1], -1.0, 1.0))
b = tf.Variable(tf.zeros([1]))

y = W * x_data + b

#Define Loss function as Mean of Squared Error
loss = tf.reduce_mean(tf.square(y - y_data))

#Create Optimizer class to minimize Losses
optimizer = tf.train.GradientDescentOptimizer(0.5)
train = optimizer.minimize(loss)

#initialize TensorFlow Variables (always)
init = tf.global_variables_initializer()
```

Execute TensorFlow Graph

#Start TensorFlow Session and carryout Variable initialization

```
sess = tf.Session()  
sess.run(init)
```

#Carryout 16 Iterations

```
for step in range(16):  
    sess.run(train)
```

#Draw Original Data

```
plt.plot(x_data, y_data, 'ro')
```


#Draw Predicted data (using calculated weight and bias after training

```
plt.plot(x_data, sess.run(W) * x_data + sess.run(b))  
plt.xlabel('x')  
plt.xlim(-2, 2)  
plt.ylim(0.1, 0.6)  
plt.ylabel('y')  
plt.legend()  
plt.show()
```


print updated weights, bias, and Loss value after current training iteration

```
print(step, sess.run(W), sess.run(b),sess.run(loss))
```


Iteration 1




Iteration 2




Iteration 3




Iteration 4




Iteration 5




Iteration 6




Iteration 7




Iteration 8




Iteration 9




Iteration 10




Iteration 11




Iteration 12




Iteration 13



Iteration 14



Iteration 15



Iteration 16

